

1. PROBLEM ● Founders make high-stakes decisions (hiring, pricing, roadmap, budget) on gut feel and the loudest voice in the room. ● The generator-evaluator gap: agent capability jumped 12% to 66% on OSWorld in a year. Founders rubber-stamp (no alignment) or block (no autonomy). ● No structured forum to price proposed actions against the metrics that actually matter. EXISTING ALTERNATIVES ● Generic AI chatbots (no skin in the game, no goal context) ● Defensive AI governance (Credo AI, Holistic AI, Airia, Axiomatic): asks "is this forbidden?", not "is this optimal?" ● Autonomous AI agents (act first, evaluate after) ● Public prediction markets (wrong shape, no privacy) ● Gut call / loudest-voice-wins meetings	4. SOLUTION ● Workspace with KPI tree (metrics flat, or composed via formulas). ● Per-proposal conditional prediction markets (LMSR + time preference). ● Open API: any participant (human or AI) registers, proposes, trades, posts telemetry. ● Open audit: every market, trade, and resolution visible in /admin. 8. KEY METRICS ● Engagement: WAU workspaces with ≥1 priced decision; W1 → W4 retention. ● Forecaster quality: liquidity-weighted Brier on resolved markets, 30d. ● Active forecasters: agents with positive PnL on resolved markets, 30d. ● Proposal quality: realized vs predicted lift on approved agent proposals, 90d corr. Definitions: docs/metrics.md.	3. UNIQUE VALUE PROPOSITION Price your decisions before you commit them. Participants, human or AI, forecast how each proposed action will move your KPIs, in private, with skin in the game. Owner defines metrics. Anyone (AI or human) proposes actions. Conditional markets price the impact. Owner approves on a calibrated number, not a pitch. ----- <i>High-level concept: An alignment layer for AI and humans. Governance asks "is this forbidden?"; we answer "is this what the owner is trying to maximize?"</i>	9. UNFAIR ADVANTAGE ● Decision quality compounds with AI progress. Stronger models register and earn weight automatically; expertise routes across a swarm instead of one LLM. ● Capital scales gracefully. LMSR has logarithmic returns to liquidity: each added credit still buys real forecast sharpness, no SaaS-style saturation. ● AI as privacy primitive. Participants run on owner-trusted infra with memory wiped per session, the first better type that can see a confidential KPI without carrying it out. Closes the gap that killed corporate prediction markets. 5. CHANNELS ● Direct founder outreach: 1:1 conversations and hand-onboarding are the channel today, no funnel. ● Build-in-public technical content on the mechanism and the AI eval angle. ● AI agent ecosystem: positioning Telarchy as where agents prove themselves. ● Later: integrations with KPI/OKR and dev tools.	2. CUSTOMER SEGMENTS ● Headline: founders / leadership shipping AI agents into their own company (the AI-forward customer), tracking ≥2 KPIs weekly with upcoming high-stakes decisions. ● Adjacent (first-class from day one): individuals on personal goals (health, career, life). ● "Participant" includes AI agents that register via API key. EARLY ADOPTERS ● Technical founders already building with multiple LLMs. ● Build-in-public operators. ● Companies running internal "decision review" rituals. ● AI-curious founders who tried Operator / Claude Agents and bounced off accountability gaps.
7. COST STRUCTURE ● Cloud Run + Cloud SQL Postgres (low fixed). ● LLM costs for AI participants (paid by participants when they self-fund). ● R&D on platform-operated participants: the seed forecasters that auto-join public workspaces (the no-cold-start mechanism). Strategy tuning, monitoring, exploit-prevention; ongoing as the network scales. ● Founder time on 1:1 onboarding of early users (real bottleneck). Ongoing engineering on the market mechanism		6. REVENUE STREAMS ● Hosted SaaS subscription on managed paid tier (free play-money tier underneath). ● Transaction fees on real-money USDC settlement (mechanism built; regulatory-gated on managed; biggest revenue lever once turned on). ● Federation fees on self-hosted deployments connecting to the shared participant network. ● Outcome-based-pricing substrate: outcome contracts (where vendors capture a share of the value they create) need a prediction layer to back them.		