

1. PROBLEM ● Founders make high-stakes decisions (hiring, pricing, roadmap, budget) on gut feel and the loudest voice in the room. ● The generator-evaluator gap: agent capability jumped 12% to 66% on OSWorld in a year. Founders rubber-stamp (no alignment) or block (no autonomy). ● No structured forum to price proposed actions against the metrics that actually matter. EXISTING ALTERNATIVES ● Generic AI chatbots (no skin in the game, no goal context) ● Defensive AI governance (Credo AI, Holistic AI, Airia, Axiomatic): asks "is this forbidden?", not "is this optimal?" ● Autonomous AI agents (act first, evaluate after) ● Public prediction markets (wrong shape, no privacy) ● Gut call / loudest-voice-wins meetings 4. SOLUTION ● Workspace with KPI tree (metrics flat, or composed via formulas). ● Per-proposal conditional prediction markets (LMSR + time preference). ● Open API: any participant (human or AI) registers, proposes, trades, posts telemetry. ● Open audit: every market, trade, and resolution visible in /admin. 8. KEY METRICS ● Wedge (now → 2026-05-27): % of cohort workspaces with ≥2 priced decisions in 4 weeks. ● Engagement (love metric): weekly active workspaces with ≥1 priced decision; W1 → W4 cohort retention. ● Forecaster quality: liquidity-weighted Brier score on resolved markets, 30d trailing. Trends down as the swarm gets smarter. ● Active forecasters: count of agents with positive cumulative PnL on resolved markets, 30d. Spam-resistant: markets are zero-sum on accuracy. ● Proposal quality: realized metric lift on approved agent-proposed proposals vs 3. UNIQUE VALUE PROPOSITION Price your decisions before you commit them. Participants, human or AI, forecast how each proposed action will move your KPIs, in private, with skin in the game. Owner defines metrics. Anyone (AI or human) proposes actions. Conditional markets price the impact. Owner approves on a calibrated number, not a pitch. ----- <i>High-level concept: An alignment layer for AI and humans. Governance asks "is this forbidden?"; we answer "is this what the owner is trying to maximize?"</i> 9. UNFAIR ADVANTAGE ● Decision quality compounds with AI progress: as stronger models register as participants, the market gives them weight automatically. Each model bets only where it has edge, so the system routes expertise across a swarm instead of betting the company on one LLM. ● Cheap-intelligence + AI-privacy: AI forecasters do not gossip, unlocking decisions humans never had a forum for. ● Working production system today (LMSR + conditional markets, composed metrics, time preference, USDC settlement) at telarchy.com. 5. CHANNELS ● Founder concierge program: direct 1:1 outreach and onboarding. The founder relationship itself is the channel for the wedge phase, no funnel. ● Build-in-public technical content on the mechanism + AI eval angle. ● AI agent ecosystem: positioning Telarchy as where agents prove themselves. ● Later: integrations with KPI/OKR and dev tools. 2. CUSTOMER SEGMENTS ● Headline: founders / leadership shipping AI agents into their own company (the AI-forward customer), tracking ≥2 KPIs weekly with upcoming high-stakes decisions. ● Adjacent (first-class from day one): individuals on personal goals (health, career, life). ● "Participant" includes AI agents that register via API key. EARLY ADOPTERS ● Technical founders already building with multiple LLMs. ● Build-in-public operators. ● Companies running internal "decision review" rituals. ● AI-curious founders who tried Operator / Claude Agents and bounced off accountability gaps. 7. COST STRUCTURE ● Cloud Run + Cloud SQL Postgres (low fixed). ● LLM costs for AI participants (paid by participants when they self-fund). ● Founder time on concierge: real bottleneck. 1:1 onboarding for the cohort across the 4-week window. ● Ongoing engineering on the matrix mechanism. 6. REVENUE STREAMS ● Hosted SaaS subscription on managed paid tier (free play-money tier underneath). ● Transaction fees on real-money USDC settlement (mechanism built; regulatory-gated on managed; biggest revenue lever once turned on). ● Federation fees on self-hosted deployments connecting to the shared participant network. ● Outcome-based-pricing substrate: outcome contracts (where vendors capture a share of the value they create) need a prediction layer to back them. <i>Today: zero revenue. The founder concierge program is the validation gate (verdict 2026-05-27).</i>
