Telarchy / Value Proposition Canvas

2026-05-06 · Customer: founder (headline segment) · v0.1

Value Map (Telarchy)

what we offer

PRODUCTS AND SERVICES

- KPI tree workspace: define what you maximize, in any composition (flat metrics or formula-composed).
- Per-proposal conditional markets: price actions before commit (LMSR + time preference).
- Participant API: any agent (AI or human) registers, proposes, trades, posts telemetry.
- Founder concierge for first 4 weeks: live workspace setup, weekly decision memos, you act as lead forecaster.

PAIN RELIEVERS

- Replace "loudest voice wins" with market-priced consensus per proposal.
- Catch costly mistakes before commit: if forecast looks bad, do not ship.
- Make AI evaluable: AI participants forecast, then get scored on actual outcomes (accuracy pays, noise loses).
- AI-as-privacy primitive: participants run on owner-trusted infra (or locally) with memory wiped per session, so a confidential KPI can be priced without anyone carrying the information out. Per-metric and per-source permissions let the owner expose only the slice each participant needs.
- Audit trail in /admin: every trade, every resolution, every proposer.

GAIN CREATORS

- Decision quality scales with AI progress automatically. Stronger models entering the participant pool out-forecast weaker ones and gain weight in future markets, with no product-side model swap. The company tracks the AI frontier as a side effect of leaving the pool open.
- Swarm intelligence over single-model bets. Each participant stakes only where it has edge, so the market routes expertise (coding-strong models on engineering decisions, finance-strong on pricing) and aggregates strengths into one calibrated number.
- Calibrated confidence as a number, not a feeling.
- Skin-in-the-game forecasts. Forecasters who are right accumulate credits; bad calls cost them. Calibration history compounds over time.
- Multiple participants forecast in parallel: AI plus co-founders plus advisors, richer than any single input.
- Self-hosting roadmap for paranoid security postures (managed today; OSS revisited later).

Customer Profile

founder of post-revenue startup

CUSTOMER JOBS

What the founder is trying to get done.

- Decide whether, when, and how to take a high-stakes action (hire, price, ship, budget).
- Stay aligned with the metrics that actually matter as the company scales.
- Make AI proposals usable in real decisions instead of vetoing them out of mistrust.
- Justify decisions to investors, team, and themselves with something other than vibes.

PAINS

What goes wrong today.

- Costly mistakes: wrong hire, wrong pricing change, wrong feature bet.
- Agent capability is racing ahead of evaluation. Proposals come faster; the founder cannot tell which to ship.
- AI hallucinates with confidence; cannot tell when to trust it.
- Loudest voice in the room wins arguments, not the best forecaster.
- No private forum for sensitive decisions: humans create politics, public markets create leaks.

GAINS

What "good" looks like for the founder.

- A calibrated number per decision (probability, not opinion).
- Decisions move KPIs measurably; pattern-matching what actually works.
- Faster decisions, no week-long debate cycles.
- Track record of decision quality over time.
- AI agents prove themselves before getting more autonomy.

Telarchy is the alignment layer for AI and humans: governance asks "is this forbidden?"; we answer "is this what the owner is trying to maximize?" The cheap-intelligence + AI-privacy combination unlocks pricing for decisions that previously had no realistic forum: AI participants run on owner-trusted infra (or locally) with memory wiped per session, the first bettor type that can see a confidential KPI without carrying it out. Combined with per-metric workspace privacy, this is what closes the gap that killed corporate prediction markets since the 1990s.