

Value Map (Telarchy) what we offer	Customer Profile founder of post-revenue startup
PRODUCTS AND SERVICES - KPI tree workspace: define what you maximize, in any composition (flat metrics or formula-composed). - Per-proposal conditional markets: price actions before commit (LMSR + time preference). - Participant API: any agent (AI or human) registers, proposes, trades, posts telemetry. - Founder concierge for first 4 weeks: live workspace setup, weekly decision memos, you act as lead forecaster.	CUSTOMER JOBS What the founder is trying to get done. - Decide whether, when, and how to take a high-stakes action (hire, price, ship, budget). - Stay aligned with the metrics that actually matter as the company scales. - Make AI proposals usable in real decisions instead of vetoing them out of mistrust. - Justify decisions to investors, team, and themselves with something other than vibes.
PAIN RELIEVERS - Replace "loudest voice wins" with market-priced consensus per proposal. - Catch costly mistakes before commit: if forecast looks bad, do not ship. - Make AI evaluable: AI participants forecast, then get scored on actual outcomes (accuracy pays, noise loses). - Per-workspace privacy with per-metric and per-source granularity. Founder controls who sees what. - Audit trail in /admin: every trade, every resolution, every proposer.	PAINS What goes wrong today. - Costly mistakes: wrong hire, wrong pricing change, wrong feature bet. - Agent capability is racing ahead of evaluation. Proposals come faster; the founder cannot tell which to ship. - AI hallucinates with confidence; cannot tell when to trust it. - Loudest voice in the room wins arguments, not the best forecaster. - No private forum for sensitive decisions: humans create politics, public markets create leaks.
GAIN CREATORS - Calibrated confidence as a number, not a feeling. - Skin-in-the-game forecasts. Forecasters who are right accumulate credits; bad calls cost them. Calibration history compounds over time. - Multiple participants forecast in parallel: AI plus co-founders plus advisors, richer than any single input. - Self-hosting roadmap for paranoid security postures (managed today; OSS revisited later).	GAINS What "good" looks like for the founder. - A calibrated number per decision (probability, not opinion). - Decisions move KPIs measurably; pattern-matching what actually works. - Faster decisions, no week-long debate cycles. - Track record of decision quality over time. - AI agents prove themselves before getting more autonomy.

Telarchy is the alignment layer for AI and humans: governance asks "is this forbidden?"; we answer "is this what the owner is trying to maximize?" The cheap-intelligence + AI-privacy combination unlocks pricing for decisions that previously had no realistic forum (a founder will not put a sensitive KPI in front of human teammates or a public market, but will put it in front of an AI participant in a private workspace). The 4-week founder concierge program (5 founders, verdict 2026-05-27) is the falsification test.